

PROBABILISTIC SYSTEMS ANALYSIS AND APPLIED PROBABILITY

MIHIR RAO

Contents

1	Probability Models and Axioms	2
1.1	Sample spaces: what can happen?	2
1.2	Events and probability axioms	3
1.3	A finite probability law: two die rolls	4
1.4	The discrete uniform law	5
1.5	The continuous uniform law	5
1.6	A countably infinite law and countable additivity	6
2	Conditioning and Bayes' Rule	8
2.1	Start with a probability model	8
2.2	Conditional probability: restrict and renormalize	8
2.3	Building a model from conditional probabilities	10
2.4	The multiplication rule	11
2.5	Total probability: combine disjoint cases	12
2.6	Bayes' rule: reverse the conditioning	13
3	Independence	14
3.1	Review: probabilities along a coin-toss tree	14
3.2	Independence of two events	15
3.3	Conditioning can change independence	16
3.4	Independence of a collection of events	18
3.5	Pairwise independence is not enough	18
3.6	The king's sibling: specify the experiment	19

LECTURE 1: PROBABILITY MODELS AND AXIOMS

Probability gives us a mathematical way to reason about uncertainty. A **PROBABILITY MODEL** specifies what can happen and assigns probabilities to events. Its two ingredients are a sample space and a probability law.

1.1 Sample spaces: what can happen?

Before assigning probabilities, decide what one complete outcome records. The detail must be enough to answer the questions of interest.

Definition 1.1 [SAMPLE SPACE]. The **SAMPLE SPACE**, denoted by Ω , is the set of all possible outcomes of an experiment. Its outcomes must be **MUTUALLY EXCLUSIVE**: only one outcome occurs in a run. They must also be **COLLECTIVELY EXHAUSTIVE**: every possible result is included.

Example: *Two rolls of a tetrahedral die.* Each roll produces one of 1, 2, 3, 4. Write X for the first result and Y for the second. One complete outcome is the ordered pair (X, Y) , so

$$\Omega = \{(x, y) : x, y \in \{1, 2, 3, 4\}\}.$$

There are $4 \times 4 = 16$ outcomes. The order matters: $(1, 2)$ and $(2, 1)$ describe different results. We can display the pairs in a grid or describe the two rolls sequentially with a tree. Each full path through the tree ends at one of the same 16 outcomes.




Figure 1: The grid and the tree describe the same sample space. A cell or a complete path records both rolls.

Example: *Two numbers in the unit interval.* If each result can be any real number from 0 to 1, the sample space is

$$\Omega = \{(x, y) : 0 \leq x \leq 1, 0 \leq y \leq 1\} = [0, 1]^2.$$

It is the entire unit square, including its boundary. The die space is finite; this space is continuous. Both still need a probability law.




Figure 2: Each point in the unit square is one possible ordered pair.

1.2 Events and probability axioms

We often ask whether a result satisfies a condition, rather than whether one particular outcome occurs.

Definition 1.2 [EVENT AND PROBABILITY LAW]. An event A is a subset of Ω : it occurs when the outcome belongs to A . A PROBABILITY LAW assigns a number $P(A)$ to each event in its domain, subject to the probability axioms below.

For a finite or countably infinite sample space, every subset can be an event. On a continuous space, the domain of the law must be chosen with care; we use ordinary geometric regions here.

Definition 1.3 [PROBABILITY AXIOMS]. For events in the domain of P :

1. NONNEGATIVITY: $P(A) \geq 0$ for every event A .
2. NORMALIZATION: $P(\Omega) = 1$.
3. ADDITIVITY: for DISJOINT events A and B , meaning $A \cap B = \emptyset$,

$$P(A \cup B) = P(A) + P(B).$$

The first two axioms say that probabilities cannot be negative and that some outcome in the sample space is certain to occur. The third lets us add probabilities when the events cannot occur together. For overlapping events, this sum would count the common outcomes twice.

Distinct individual outcomes give disjoint singleton events. Repeatedly applying additivity therefore gives

$$P(\{s_1, \dots, s_k\}) = P(\{s_1\}) + \dots + P(\{s_k\}).$$

We may abbreviate $P(\{s\})$ to $P(s)$. The two-event axiom handles finite unions; an infinite union will require a stronger form of additivity.

1.3 A finite probability law: two die rolls

Return to the 16 ordered die outcomes and assign probability $1/16$ to each. The probabilities are nonnegative and sum to $16(1/16) = 1$. An event containing k outcomes has probability $k/16$.

Example: *Either (1, 1) or (1, 2).* These are two distinct outcomes, so

$$P(\{(1, 1), (1, 2)\}) = \frac{1}{16} + \frac{1}{16} = \frac{1}{8}.$$

Example: *The first roll is 1.* The second roll can be 1, 2, 3, or 4. Thus

$$\{X = 1\} = \{(1, 1), (1, 2), (1, 3), (1, 4)\}, \quad P(X = 1) = \frac{4}{16} = \frac{1}{4}.$$

The notation $\{X = 1\}$ names the subset of outcomes satisfying the condition.

Example: *The sum is odd.* One roll must be odd and the other even. There are two odd values and two even values, so there are $2 \times 2 = 4$ pairs of each order. Hence

$$P(X + Y \text{ is odd}) = \frac{4 + 4}{16} = \frac{1}{2}.$$

Example: *The smaller roll is 2.* Both rolls must be at least 2, and at least one must equal 2. The event is

$$\{\min(X, Y) = 2\} = \{(2, 2), (2, 3), (2, 4), (3, 2), (4, 2)\}.$$

There are five outcomes, giving

$$P(\min(X, Y) = 2) = \frac{5}{16}.$$

Counting three pairs with $X = 2$ and three with $Y = 2$ would count $(2, 2)$ twice, so it must be included only once.




Figure 3: Highlighted cells show the last three die events. Each cell has probability $1/16$. In the last panel, the corner $(2, 2)$ is counted once.

1.4 The discrete uniform law

When a finite sample space has N equally likely outcomes, normalization forces each outcome to have probability $1/N$. This is the **DISCRETE UNIFORM LAW**. Additivity then gives

$$P(A) = \frac{|A|}{|\Omega|},$$

where $|A|$ is the number of outcomes in A and $|\Omega| = N$. Computing a probability becomes a counting problem. Fair coins, fair dice, and well-shuffled decks are modeled this way. The formula relies on the outcomes being equally likely.

1.5 The continuous uniform law

Now give the unit square a **CONTINUOUS UNIFORM LAW**: the probability of a region is its area. The square's total area is 1, so the law is normalized. Equal-area regions have equal probabilities, and areas add for disjoint regions.

Example: *The two numbers sum to at most $1/2$.* The event is

$$A = \{(x, y) \in [0, 1]^2 : x + y \leq \frac{1}{2}\}.$$

It is the triangle bounded by the two coordinate axes and the line $x + y = 1/2$. Both perpendicular sides have length $1/2$, so

$$P(X + Y \leq \frac{1}{2}) = \text{area}(A) = \frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{8}.$$




Figure 4: The shaded triangle has area $1/8$. The point $(0.5, 0.3)$ is a possible outcome, but its singleton has zero area.

Example: *The exact pair* $(0.5, 0.3)$. A single point has zero area, so

$$P((X, Y) = (0.5, 0.3)) = 0.$$

The point is still in the sample space. A probability-zero event need not be impossible: in this model every individual point has probability zero, while the whole square has probability one.

1.6 A countably infinite law and countable additivity

Consider the sample space $\Omega = \{1, 2, 3, \dots\}$ with

$$P(n) = 2^{-n}, \quad n = 1, 2, \dots$$

Each probability is positive, and the geometric series checks normalization:

$$\sum_{n=1}^{\infty} 2^{-n} = \frac{1/2}{1 - 1/2} = 1.$$




Figure 5: The first four outcome probabilities. Each is half the preceding one, and the pattern continues for every positive integer.

Example: *The outcome is even.* The event $\{2, 4, 6, \dots\}$ contains infinitely many disjoint singleton events. Adding their probabilities gives

$$\begin{aligned} P(\{2, 4, 6, \dots\}) &= \sum_{k=1}^{\infty} 2^{-2k} = \frac{1}{4} + \frac{1}{16} + \frac{1}{64} + \dots \\ &= \frac{1/4}{1 - 1/4} = \frac{1}{3}. \end{aligned}$$

This calculation uses an infinite sum, which finite additivity alone does not justify.

Definition 1.4 [COUNTABLE ADDITIVITY]. If $A_1, A_2, \dots$ are pairwise disjoint events, the COUNTABLE ADDITIVITY axiom requires

$$P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(A_i).$$

This is the full additivity axiom for a probability law. It covers a sequence indexed by the positive integers and includes finite additivity as a special case. It justifies the even-outcome calculation above. It does not permit summing singleton probabilities over all the points of the unit square: that collection is uncountable.

Revision. A model needs both a sample space and a probability law. Events are sets of outcomes. Count outcomes under a finite uniform law, measure area under the uniform law on the unit square, and use countable additivity for a countable union of disjoint events.

Common mistakes: using the counting formula without equal likelihood, counting a shared outcome twice, or treating probability zero as impossibility.

Check yourself: Why do two tetrahedral-die rolls give 16 outcomes? Why does $\min(X, Y) = 2$ select five of them? Why can an infinite sum give the probability of an even integer, but cannot add up the probabilities of every point in the unit square?

Source. John Tsitsiklis, MIT 6.041/6.431, *Probabilistic Systems Analysis and Applied Probability*, Fall 2010, [Lecture 1: Probability Models and Axioms](#), PDF pp. 1–3. The calculations are completed and the diagrams redrawn here. Adapted from MIT OpenCourseWare under [CC BY-NC-SA 4.0](#).

LECTURE 2: CONDITIONING AND BAYES' RULE

Learning that an event occurred changes which outcomes remain possible and how their probabilities should be compared. Conditioning makes this update precise. It also gives three useful tools: multiplying probabilities along a sequence, combining disjoint cases, and reversing a conditional probability.

2.1 Start with a probability model

The sample space Ω lists mutually exclusive, collectively exhaustive outcomes at the detail needed for the question. Events are subsets of Ω . Their probabilities are nonnegative, satisfy $P(\Omega) = 1$, and add over disjoint events. For a disjoint sequence $A_1, A_2, \dots$,

$$P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(A_i).$$

To solve a problem, specify the sample space, define its probability law, identify the event of interest, and then calculate. A calculation is only meaningful once these ingredients are clear.

2.2 Conditional probability: restrict and renormalize

Suppose we learn that B occurred. Outcomes outside B can be discarded. Among the outcomes in B , the event A occurs exactly on $A \cap B$. We compare the probability of this intersection with the probability of the new universe B .

Definition 2.1 [CONDITIONAL PROBABILITY]. For events A and B with $P(B) > 0$, the **CONDITIONAL PROBABILITY** of A given B is

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}.$$

This definition does not assign a value when $P(B) = 0$.

Restricting to B keeps only the outcomes compatible with the information. Dividing by $P(B)$ rescales the remaining probability to one: $P(B \mid B) = P(B)/P(B) = 1$.




Figure 6: Conditioning keeps B and compares $A \cap B$ with all of B . This is a schematic set diagram; region sizes do not represent probabilities.

Example: *The minimum of two die rolls is 2.* Let X and Y be the two results of a tetrahedral die, with each of the 16 ordered pairs assigned probability $1/16$. Suppose we know

$$B = \{\min(X, Y) = 2\} = \{(2, 2), (2, 3), (2, 4), (3, 2), (4, 2)\}.$$

Then $P(B) = 5/16$. Each remaining pair has conditional probability $(1/16)/(5/16) = 1/5$.

Let $M = \max(X, Y)$. The event $M = 1$ has no outcome in common with B : if the minimum is 2, the maximum cannot be 1. Therefore

$$P(M = 1 \mid B) = \frac{0}{5/16} = 0.$$

For $M = 2$, the only compatible pair is $(2, 2)$, giving

$$P(M = 2 \mid B) = \frac{1/16}{5/16} = \frac{1}{5}.$$

The denominator counts the five outcomes allowed by the information, not all 16 outcomes in the original space.




Figure 7: The five colored cells form B . The darker cell $(2, 2)$ is the only one where $M = 2$. Each colored cell has probability $1/5$ given B .

2.3 Building a model from conditional probabilities

A model can be specified in stages. First give the probability of one event, then give the probabilities of later events under each possible earlier result. Write A^c for the COMPLEMENT of A , the event that A does not occur.

Example: *An airplane and a radar reading.* Let A mean that an airplane is flying above, and let B mean that something registers on the radar screen. The model is

$$P(A) = 0.05, \quad P(A^c) = 0.95,$$

with the following conditional probabilities:

$$\begin{aligned} P(B | A) &= 0.99 & P(B^c | A) &= 0.01, \\ P(B | A^c) &= 0.10 & P(B^c | A^c) &= 0.90. \end{aligned}$$

Thus a registration is likely when a plane is present, but registrations also occur when no plane is present. Within each condition, the two radar probabilities sum to one.

Rearranging the conditional-probability definition gives

$$P(A \cap B) = P(A)P(B | A).$$

Multiply the probability of the first stage by the conditional probability of the second stage to obtain the probability of a complete path.




Figure 8: Each path specifies both plane presence and radar status. Multiply the two factors on a path to obtain the joint probability on the right.

For a plane and a registration,

$$P(A \cap B) = 0.05(0.99) = 0.0495.$$

A registration can come from either $A \cap B$ or $A^c \cap B$. These events are disjoint, so

$$\begin{aligned} P(B) &= P(A \cap B) + P(A^c \cap B) \\ &= 0.05(0.99) + 0.95(0.10) = 0.1445. \end{aligned}$$

Once a registration is observed, the probability of a plane is

$$P(A | B) = \frac{P(A \cap B)}{P(B)} = \frac{0.0495}{0.1445} = \frac{99}{289} \approx 0.3426.$$

The two directions answer different questions: $P(B | A) = 0.99$ describes the radar given a plane; $P(A | B) \approx 0.3426$ describes plane presence given a registration. The latter must account for how often planes are present and how often the radar registers without one.

2.4 The multiplication rule

For three events, a complete path specifies A , then B given A , then C given both earlier events. Each factor must use all the information accumulated so far.

Proposition 2.2 [MULTIPLICATION RULE]. *If $P(A) > 0$ and $P(A \cap B) > 0$, then*

$$P(A \cap B \cap C) = P(A)P(B | A)P(C | A \cap B).$$

The positivity assumptions ensure that the conditional probabilities on the right are defined.

Proof. Apply the conditional-probability definition twice: first condition on the whole event $A \cap B$, then express $P(A \cap B)$ by conditioning on A .

$$\begin{aligned} P(A \cap B \cap C) &= P(A \cap B)P(C | A \cap B) \\ &= P(A)P(B | A)P(C | A \cap B). \end{aligned}$$

□




Figure 9: One path through a three-stage model. The last factor conditions on $A \cap B$, because both earlier events have been specified.

2.5 Total probability: combine disjoint cases

Sometimes the probability of B is easier to calculate after dividing the sample space into cases. Compute the contribution from each case, then add the contributions.

Definition 2.3 [PARTITION]. Events A_1, A_2, A_3 form a PARTITION of Ω if they are pairwise disjoint and cover the whole space:

$$A_i \cap A_j = \emptyset \quad (i \neq j), \quad A_1 \cup A_2 \cup A_3 = \Omega.$$

Every outcome belongs to exactly one of the three cases.

Theorem 2.4 [TOTAL PROBABILITY]. For a partition A_1, A_2, A_3 with $P(A_i) > 0$ for each i ,

$$P(B) = P(A_1)P(B \mid A_1) + P(A_2)P(B \mid A_2) + P(A_3)P(B \mid A_3).$$

Proof. Each outcome of B belongs to exactly one part of the partition, so $B = (A_1 \cap B) \cup (A_2 \cap B) \cup (A_3 \cap B)$ is a disjoint union. Additivity gives

$$P(B) = P(A_1 \cap B) + P(A_2 \cap B) + P(A_3 \cap B).$$

For each piece, use $P(A_i \cap B) = P(A_i)P(B \mid A_i)$. Substituting these three products gives the formula. $\square$




Figure 10: The partition cuts B into three disjoint pieces. Their probabilities add to $P(B)$. The geometry is schematic.

Each term weights a conditional probability by the probability of its case. In the radar example, the two cases are A and A^c , giving the calculation $P(B) = 0.05(0.99) + 0.95(0.10)$.

2.6 Bayes' rule: reverse the conditioning

Suppose we know how likely B is under each case A_i , but observe B and want to infer which case occurred. The PRIOR PROBABILITIES $P(A_i)$ describe the cases before using this observation. The updated probabilities are $P(A_i | B)$.

Theorem 2.5 [BAYES' RULE]. *For the partition above, if $P(A_i) > 0$ for each i and $P(B) > 0$, then*

$$P(A_i | B) = \frac{P(A_i)P(B | A_i)}{\sum_{j=1}^3 P(A_j)P(B | A_j)}, \quad i = 1, 2, 3.$$

Proof. The conditional-probability definition gives

$$P(A_i | B) = \frac{P(A_i \cap B)}{P(B)}.$$

The multiplication rule supplies the numerator, $P(A_i \cap B) = P(A_i)P(B | A_i)$. The total probability theorem supplies the denominator by adding the contributions from all three cases. Substituting both expressions yields Bayes' rule. $\square$

In [figure 10](#), the numerator is the probability of the piece $A_i \cap B$ and the denominator is the probability of all of B . Thus Bayes' rule compares one case's contribution with the total probability of the observation. In the radar model, it updates $P(A) = 0.05$ to $P(A | B) \approx 0.3426$ after a registration.

Revision. Conditioning restricts the possible outcomes and divides by the probability of the conditioning event. Multiply conditional factors along a path; add disjoint paths for total probability. Bayes' rule divides one case's contribution by the total probability of the observed event.

Common mistakes: interchanging $P(A | B)$ and $P(B | A)$, using a denominator that ignores the information, or forgetting earlier events in the last factor of the multiplication rule.

Check yourself: Why is $P(M = 2 | \min(X, Y) = 2) = 1/5$? Why is the radar model's $P(B) = 0.1445$? Which two ingredients, besides $P(B | A)$, are needed to calculate $P(A | B)$?

Source. John Tsitsiklis, MIT 6.041/6.431, *Probabilistic Systems Analysis and Applied Probability*, Fall 2010, [Lecture 2: Conditioning and Bayes' Rule](#), PDF pp. 1–2. The supplied calculations and derivations are completed, and the diagrams redrawn here. Adapted from MIT OpenCourseWare under [CC BY-NC-SA 4.0](#).

LECTURE 3: INDEPENDENCE

Independent events do not change one another's probabilities. Independence depends on the probability law: conditioning can remove it or create it.

3.1 Review: probabilities along a coin-toss tree

For $P(B) > 0$, conditioning gives $P(A \mid B) = P(A \cap B)/P(B)$. The multiplication rule, total probability, and Bayes' rule then give

$$\begin{aligned} P(A \cap B) &= P(A)P(B \mid A) = P(B)P(A \mid B), \\ P(B) &= P(A)P(B \mid A) + P(A^c)P(B \mid A^c), \\ P(A \mid B) &= \frac{P(A)P(B \mid A)}{P(B)}. \end{aligned}$$

Each conditional factor requires a conditioning event of positive probability. Multiply along a path and add disjoint paths.

Example: *Three tosses of a biased coin.* At each stage, the probability of heads is p and of tails is $1 - p$, regardless of earlier results. The sample space consists of the eight strings $HHH, HHT, HTH, HTT, THH, THT, TTH, TTT$.




Figure 11: Each box gives the accumulated result and the probability of the newest toss. Multiply the three factors on a path. These factors remain p or $1 - p$ after every earlier history.

The path THT has probability

$$P(THT) = (1 - p)p(1 - p) = p(1 - p)^2.$$

Exactly one head occurs on the three disjoint paths HTT, THT, TTH . Each has the same probability, so

$$P(\text{exactly one head}) = 3p(1 - p)^2.$$

Given exactly one head, the first toss is heads only on HTT . For $0 < p < 1$,

$$P(\text{first toss is } H \mid \text{exactly one head}) = \frac{p(1 - p)^2}{3p(1 - p)^2} = \frac{1}{3}.$$

The bias cancels because all three remaining paths have equal probability. At $p = 0$ or $p = 1$, the conditioning event has probability zero, so this conditional probability is undefined.

3.2 Independence of two events

If learning that A occurred gives no information about B , then $P(B \mid A) = P(B)$, provided $P(A) > 0$. Substitution in the multiplication rule gives a definition that also works for probability-zero events.

Definition 3.1 [INDEPENDENCE]. Events A and B are INDEPENDENT if

$$P(A \cap B) = P(A)P(B).$$

The definition is symmetric in A and B . When $P(A) > 0$, it is equivalent to $P(B \mid A) = P(B)$; when $P(B) > 0$, it is equivalent to $P(A \mid B) = P(A)$.

Proof. If $P(A) > 0$, divide $P(A \cap B) = P(A)P(B)$ by $P(A)$ to obtain $P(B \mid A) = P(B)$. Conversely, multiplying that conditional equality by $P(A)$ gives the defining product. Interchanging A and B proves the other equivalence. $\square$

If $P(A) = 0$, then $A \cap B \subseteq A$ also has probability zero. Thus $P(A \cap B) = 0 = P(A)P(B)$ for every B , even though conditioning on A is undefined. Independence is a numerical property of the probability law; the shape of a set diagram alone does not establish it.

3.3 Conditioning can change independence

After learning that C occurred, use the conditional probability law $P(\cdot | C)$ throughout. Independence in this new universe is tested using conditional probabilities.

Definition 3.2 [CONDITIONAL INDEPENDENCE]. For $P(C) > 0$, events A and B are **CONDITIONALLY INDEPENDENT** given C if

$$P(A \cap B | C) = P(A | C)P(B | C).$$




Figure 12: Within C , the pieces $A \cap C$ and $B \cap C$ are disjoint. The diagram is schematic; areas do not specify probabilities.

Suppose A and B are independent before conditioning, but C meets each of them while avoiding $A \cap B$. If both $P(A | C)$ and $P(B | C)$ are positive, then

$$P(A \cap B | C) = 0 \quad \text{while} \quad P(A | C)P(B | C) > 0.$$

They are therefore dependent given C . Unconditional independence does not guarantee conditional independence.

Example: *A coin chosen once, then tossed repeatedly.* Choose Coin A or Coin B with probability $1/2$ each. Coin A has head probability 0.9 ; Coin B has head probability 0.1 . Toss the chosen coin repeatedly, with independent tosses conditional on its identity. Let G be the event that Coin A was chosen, and let H_i mean that toss i is heads. Then

$$P(H_i | G) = 0.9, \quad P(H_i | G^c) = 0.1.$$

Knowing the identity fixes the head probability. Within either conditional model, earlier tosses give no further information about later tosses.




Figure 13: The first choice selects one coin for all tosses. The boxes give the probability of each new step conditional on the preceding history. The two tosses shown are independent given the chosen coin.

Without knowing the coin, total probability gives

$$P(H_{11}) = \frac{1}{2}(0.9) + \frac{1}{2}(0.1) = \frac{1}{2}.$$

Now suppose the first ten tosses were heads. Write $D = H_1 \cap \dots \cap H_{10}$. Conditional independence within each coin model gives $P(D \mid G) = 0.9^{10}$ and $P(D \mid G^c) = 0.1^{10}$. The observation is far more likely with Coin A. Bayes' rule gives

$$P(G \mid D) = \frac{0.9^{10}}{0.9^{10} + 0.1^{10}}.$$

For the next toss, divide the probability of eleven heads by that of ten:

$$\begin{aligned} P(H_{11} \mid D) &= \frac{\frac{1}{2}(0.9^{11}) + \frac{1}{2}(0.1^{11})}{\frac{1}{2}(0.9^{10}) + \frac{1}{2}(0.1^{10})} \\ &= \frac{0.9^{11} + 0.1^{11}}{0.9^{10} + 0.1^{10}} \approx 0.9. \end{aligned}$$

This differs from $P(H_{11}) = 1/2$. Earlier results reveal information about the shared coin and change the probability of later results. Thus conditional independence does not guarantee unconditional independence.

3.4 Independence of a collection of events

For a collection, the idea is stronger than checking two events at a time: information about any group of events must give no information about the remaining events.

Definition 3.3 [INDEPENDENCE OF A COLLECTION]. Events $A_1, \dots, A_n$ are **MUTUALLY INDEPENDENT** if, for every nonempty subset $S \subseteq \{1, \dots, n\}$,

$$P\left(\bigcap_{i \in S} A_i\right) = \prod_{i \in S} P(A_i).$$

Here $\prod$ means multiplication over the indicated indices. The condition must hold for every subcollection, not just all n events together.

For three events, this requires the three pairwise products and also

$$P(A_1 \cap A_2 \cap A_3) = P(A_1)P(A_2)P(A_3).$$

For a larger independent collection, even combinations of events and their complements involving disjoint groups of indices give no information about one another. For example, when $n \geq 6$ and $P(A_5 \cap A_6^c) > 0$,

$$P(A_1 \cap (A_2^c \cup A_3) \mid A_5 \cap A_6^c) = P(A_1 \cap (A_2^c \cup A_3)).$$

The condition concerns indices 5, 6, while the event of interest concerns indices 1, 2, 3.

3.5 Pairwise independence is not enough

Events are **PAIRWISE INDEPENDENT** if each pair of distinct events is independent. This checks only subcollections of size two. It need not establish mutual independence.

Example: *Two fair coin tosses and whether they match.* Consider two independent fair tosses, with the four equally likely outcomes HH, HT, TH, TT . Define

$$A = \{HH, HT\}, \quad B = \{HH, TH\}, \quad C = \{HH, TT\}.$$

Thus A means the first toss is heads, B the second is heads, and C that both results are the same.

		Second toss		
		H	T	
First toss	H	HH	HT	A: first toss is H
	T	TH	TT	B: second toss is H
				C: the tosses match

Figure 14: Every cell has probability $1/4$. The two colored cells form C . Knowing both A and B identifies HH , which is in C .

Each event contains two outcomes, so $P(A) = P(B) = P(C) = 1/2$. Each pairwise intersection is the singleton $\{HH\}$:

$$P(A \cap B) = P(A \cap C) = P(B \cap C) = \frac{1}{4} = \frac{1}{2} \cdot \frac{1}{2}.$$

All three pairs are independent. But the triple intersection is also $\{HH\}$, giving

$$P(A \cap B \cap C) = \frac{1}{4} \neq \frac{1}{8} = P(A)P(B)P(C).$$

Equivalently,

$$P(C \mid A \cap B) = \frac{1/4}{1/4} = 1 \neq \frac{1}{2} = P(C).$$

Neither toss alone changes the probability of matching; knowing both tosses determines whether they match. Pairwise independence misses this dependence on combined information.

3.6 The king's sibling: specify the experiment

Example: A king comes from a family with two children. What is the probability that his sibling is female? Assume the parents decided in advance to have exactly two children. Each child is a boy or girl with probability $1/2$, independently of the other. Assume that a boy takes precedence in succession, so a king exists exactly when at least one child is a boy.

Order the children by birth. The four outcomes BB, BG, GB, GG each have probability $1/4$, where B means boy and G girl. The information that there is a king is the event

$$K = \{BB, BG, GB\}, \quad P(K) = \frac{3}{4}.$$

His sibling is female in BG or GB . Writing F for this event within K gives

$$P(F \mid K) = \frac{P(F \cap K)}{P(K)} = \frac{2/4}{3/4} = \frac{2}{3}.$$

		Younger child		
		B	G	
Older child	B	BB	BG	
	G	GB	GG excluded	

K : at least one boy
 $F \cap K$: BG or GB
 $P(F | K) = 2/3$

Figure 15: Under the fixed two-child model, conditioning on K removes GG . The other three outcomes are equally likely; two give a female sibling. The darker cells are $F \cap K$.

The answer depends on the experiment. If the parents instead have children until the first boy and then stop, having exactly two children forces the outcome GB . Under that rule, the king's sibling is female with probability one. The phrase "a family with two children" alone does not specify how the family came to have two children.

Revision. Independence is the product condition $P(A \cap B) = P(A)P(B)$. Conditional independence uses the same condition under $P(\cdot | C)$. Mutual independence checks every subcollection; pairwise independence checks only pairs.

Common mistakes: carrying independence into a conditional model without checking it, inferring mutual independence from pairwise products, or leaving the experiment's selection and stopping rules unspecified.

Check yourself: Why does exactly one head give a first-toss head probability of $1/3$, regardless of the bias? Why do ten heads change the next-toss probability for the unknown coin? Which product fails in the matching-toss example? What assumptions produce the sibling probability $2/3$?

Source. John Tsitsiklis, MIT 6.041/6.431, *Probabilistic Systems Analysis and Applied Probability*, Fall 2010, [Lecture 3: Independence](#), PDF pp. 1–2, and the [accompanying transcript](#) (conditional independence and the sibling model, especially 41:29–46:28). The supplied calculations and derivations are completed, and the diagrams redrawn here. Adapted from MIT OpenCourseWare under [CC BY-NC-SA 4.0](#).